

## 6.Problem: Assuming a set of documents that need to be classified, use the naïve Bayesian Classifier model to perform this task. Built-in Libraries can be used to write the program. Calculate the accuracy, precision, and recall for your data set.

### Example1

Loading the 20 newsgroups dataset : The dataset is called “Twenty Newsgroups”. Here is the official description, quoted from the website:<http://qwone.com/~jason/20Newsgroups/>

The 20 Newsgroups data set is a collection of approximately 20,000 newsgroup documents, partitioned (nearly) evenly across 20 different newsgroups. To the best of our knowledge, it was originally collected by Ken Lang, probably for his paper “Newsweeder: Learning to filter netnews,” though he does not explicitly mention this collection. The 20 newsgroups collection has become a popular data set for experiments in text applications of machine learning techniques, such as text classification and text clustering.

In [1]:

```
# Author : Dr.Thyagaraju G S , Context Innovations Lab , DEpt of CSE , SDMIT - Ujire
# Date : July 11 2018
from sklearn.datasets import fetch_20newsgroups
twenty_train = fetch_20newsgroups(subset='train', shuffle=True)
x = len(twenty_train.target_names)
print("\n The number of categories:",x)
print("\n The %d Different Categories of 20Newsgroups\n" %x)
i=1
for cat in twenty_train.target_names:
    print("Category[%d]:" %i,cat)
    i=i+1
print("\n Length of training data is",len(twenty_train.data))
print("\n Length of file names is ",len(twenty_train.filenames))

print("\n The Content/Data of First File is :\n")

print(twenty_train.data[0])
```

Downloading 20news dataset. This may take a few minutes.  
Downloading dataset from <https://ndownloader.figshare.com/files/5975967> (14 MB)

The number of categories: 20

The 20 Different Categories of 20Newsgroups

```
Category[1]: alt.atheism
Category[2]: comp.graphics
Category[3]: comp.os.ms-windows.misc
Category[4]: comp.sys.ibm.pc.hardware
Category[5]: comp.sys.mac.hardware
Category[6]: comp.windows.x
Category[7]: misc.forsale
Category[8]: rec.autos
Category[9]: rec.motorcycles
Category[10]: rec.sport.baseball
Category[11]: rec.sport.hockey
Category[12]: sci.crypt
Category[13]: sci.electronics
Category[14]: sci.med
Category[15]: sci.space
Category[16]: soc.religion.christian
Category[17]: talk.politics.guns
Category[18]: talk.politics.mideast
Category[19]: talk.politics.misc
Category[20]: talk.religion.misc
```

Length of training data is 11314

Length of file names is 11314

The Content/Data of First File is :

From: lerxst@wam.umd.edu (where's my thing)  
Subject: WHAT car is this!?  
Nntp-Posting-Host: rac3.wam.umd.edu  
Organization: University of Maryland, College Park  
Lines: 15

I was wondering if anyone out there could enlighten me on this car I saw the other day. It was a 2-door sports car, looked to be from the late 60s/early 70s. It was called a Bricklin. The doors were really small. In addition, the front bumper was separate from the rest of the body. This is all I know. If anyone can tell me a model name, engine specs, years of production, where this car is made, history, or whatever info you have on this funky looking car, please e-mail.

Thanks,

- IL

---- brought to you by your neighborhood Lerxst ----

In [2]:

```
print("\n The Contents/Data of First 10 Files is in Training Data :\n")  
  
for i in range(0,10):  
    print("\n FILE NO:%d \n"%(i+1))  
    print(twenty_train.data[i])
```

The Contents/Data of First 10 Files is in Training Data :

FILE NO:1

From: lerxst@wam.umd.edu (where's my thing)  
Subject: WHAT car is this!?  
Nntp-Posting-Host: rac3.wam.umd.edu  
Organization: University of Maryland, College Park  
Lines: 15

I was wondering if anyone out there could enlighten me on this car I saw the other day. It was a 2-door sports car, looked to be from the late 60s/early 70s. It was called a Bricklin. The doors were really small. In addition, the front bumper was separate from the rest of the body. This is all I know. If anyone can tell me a model name, engine specs, years of production, where this car is made, history, or whatever info you have on this funky looking car, please e-mail.

Thanks,

- IL

---- brought to you by your neighborhood Lerxst ----

FILE NO:2

From: guykuo@carson.u.washington.edu (Guy Kuo)  
Subject: SI Clock Poll - Final Call  
Summary: Final call for SI clock reports  
Keywords: SI, acceleration, clock, upgrade  
Article-I.D.: shelley.1qvfo9INNC3s  
Organization: University of Washington  
Lines: 11  
NNTP-Posting-Host: carson.u.washington.edu

A fair number of brave souls who upgraded their SI clock oscillator have shared their experiences for this poll. Please send a brief message detailing your experiences with the procedure. Top speed attained, CPU rated speed,

add on cards and adapters, heat sinks, hour of usage per day, floppy disk functionality with 800 and 1.4 m floppies are especially requested.

I will be summarizing in the next two days, so please add to the network knowledge base if you have done the clock upgrade and haven't answered this poll. Thanks.

Guy Kuo <guykuo@u.washington.edu>

FILE NO:3

From: twillis@ecn.purdue.edu (Thomas E Willis)  
Subject: PB questions...  
Organization: Purdue University Engineering Computer Network  
Distribution: usa  
Lines: 36

well folks, my mac plus finally gave up the ghost this weekend after starting life as a 512k way back in 1985. sooo, i'm in the market for a new machine a bit sooner than i intended to be...

i'm looking into picking up a powerbook 160 or maybe 180 and have a bunch of questions that (hopefully) somebody can answer:

\* does anybody know any dirt on when the next round of powerbook introductions are expected? i'd heard the 185c was supposed to make an appearance "this summer" but haven't heard anymore on it - and since i don't have access to maclean, i was wondering if anybody out there had more info...

\* has anybody heard rumors about price drops to the powerbook line like the ones the duo's just went through recently?

\* what's the impression of the display on the 180? i could probably swing a 180 if i got the 80Mb disk rather than the 120, but i don't really have a feel for how much "better" the display is (yea, it looks great in the store, but is that all "wow" or is it really that good?). could i solicit some opinions of people who use the 160 and 180 day-to-day on if its worth taking the disk size and money hit to get the active display? (i realize this is a real subjective question, but i've only played around with the machines in a computer store briefly and figured the opinions of somebody who actually uses the machine daily might prove helpful).

\* how well does hellcats perform? ;)

thanks a bunch in advance for any info - if you could email, i'll post a summary (news reading time is at a premium with finals just around the corner... :( )

--

Tom Willis \ twillis@ecn.purdue.edu \ Purdue Electrical Engineering

-----  
"Convictions are more dangerous enemies of truth than lies." - F. W. Nietzsche

FILE NO:4

From: jgreen@amber (Joe Green)  
Subject: Re: Weitek P9000 ?  
Organization: Harris Computer Systems Division  
Lines: 14  
Distribution: world  
NNTP-Posting-Host: amber.ssd.csd.harris.com  
X-Newsreader: TIN [version 1.1 PL9]

Robert J.C. Kyanko (rob@rjck.UUCP) wrote:  
> abraxis@iastate.edu writes in article <abraxis.734340159@class1.iastate.edu>:  
> > Anyone know about the Weitek P9000 graphics chip?  
> As far as the low-level stuff goes, it looks pretty nice. It's got this  
> quadrilateral fill command that requires just the four points.

Do you have Weitek's address/phone number? I'd like to get some information about this chip.

--

Joe Green Harris Corporation

cc: green@harris.com  
jgreen@csd.harris.com Computer Systems Division  
"The only thing that really scares me is a person with no sense of humor."  
-- Jonathan Winters

FILE NO:5

From: jcm@head-cfa.harvard.edu (Jonathan McDowell)  
Subject: Re: Shuttle Launch Question  
Organization: Smithsonian Astrophysical Observatory, Cambridge, MA, USA  
Distribution: sci  
Lines: 23

From article <C5owCB.n3p@world.std.com>, by tombaker@world.std.com (Tom A Baker):  
>>In article <C5JLwx.4H9.1@cs.cmu.edu>, ETRAT@ttacs1.ttu.edu (Pack Rat) writes...  
>>>"Clear caution & warning memory. Verify no unexpected  
>>>errors. ...". I am wondering what an "expected error" might  
>>>be. Sorry if this is a really dumb question, but  
>  
> Parity errors in memory or previously known conditions that were waived.  
> "Yes that is an error, but we already knew about it"  
> I'd be curious as to what the real meaning of the quote is.  
>  
> tom

My understanding is that the 'expected errors' are basically known bugs in the warning system software - things are checked that don't have the right values in yet because they aren't set till after launch, and suchlike. Rather than fix the code and possibly introduce new bugs, they just tell the crew 'ok, if you see a warning no. 213 before liftoff, ignore it'.

- Jonathan

FILE NO:6

From: dfo@vttoulu.tko.vtt.fi (Foxvog Douglas)  
Subject: Re: Rewording the Second Amendment (ideas)  
Organization: VTT  
Lines: 58

In article <lrleul\$4t@transfer.stratus.com> cdt@sw.stratus.com (C. D. Tavares) writes:  
>In article <1993Apr20.083057.16899@ousrvr.oulu.fi>, dfo@vttoulu.tko.vtt.fi (Foxvog Douglas) writes:  
>> In article <lqv87v\$4j3@transfer.stratus.com> cdt@sw.stratus.com (C. D. Tavares) writes:  
>> >In article <C5n3GI.F8F@ulowell.ulowell.edu>, jrutledg@cs.ulowell.edu (John Lawrence Rutledge) writes:  
>  
>> >> The massive destructive power of many modern weapons, makes the  
>> >> cost of an accidental or criminal usage of these weapons to great.  
>> >> The weapons of mass destruction need to be in the control of  
>> >> the government only. Individual access would result in the  
>> >> needless deaths of millions. This makes the right of the people  
>> >> to keep and bear many modern weapons non-existent.  
  
>> >Thanks for stating where you're coming from. Needless to say, I  
>> >disagree on every count.  
  
>> You believe that individuals should have the right to own weapons of  
>> mass destruction? I find it hard to believe that you would support a  
>> neighbor's right to keep nuclear weapons, biological weapons, and nerve  
>> gas on his/her property.  
  
>> If we cannot even agree on keeping weapons of mass destruction out of  
>> the hands of individuals, can there be any hope for us?  
  
>I don't sign any blank checks.

Of course. The term must be rigidly defined in any bill.

>When Doug Foxvog says "weapons of mass destruction," he means CBW and  
>nukes. When Sarah Brady says "weapons of mass destruction" she means  
>Street Sweeper shotguns and semi-automatic SKS rifles.

rocket sweeper shotgun and some automatics she likes.

I doubt she uses this term for that. You are using a quote allegedly from her, can you back it up?

>When John

>Lawrence Rutledge says "weapons of mass destruction," and then immediately  
>follows it with:

>>> The US has thousands of people killed each year by handguns,  
>>> this number can easily be reduced by putting reasonable restrictions  
>>> on them.

>...what does Rutledge mean by the term?

I read the article as presenting first an argument about weapons of mass destruction (as commonly understood) and then switching to other topics. The first point evidently was to show that not all weapons should be allowed, and then the later analysis was, given this understanding, to consider another class.

>cdt@rocket.sw.stratus.com --If you believe that I speak for my company,  
>OR cdt@vos.stratus.com write today for my special Investors' Packet...

--

doug foxvog  
douglas.foxvog@vtt.fi

FILE NO:7

From: bmdelane@quads.uchicago.edu (brian manning delaney)  
Subject: Brain Tumor Treatment (thanks)  
Reply-To: bmdelane@midway.uchicago.edu  
Organization: University of Chicago  
Lines: 12

There were a few people who responded to my request for info on treatment for astrocytomas through email, whom I couldn't thank directly because of mail-bouncing probs (Sean, Debra, and Sharon). So I thought I'd publicly thank everyone.

Thanks!

(I'm sure glad I accidentally hit "rn" instead of "rm" when I was trying to delete a file last September. "Hmmm... 'News?' What's this?")....)

-Brian

FILE NO:8

From: bgrubb@dante.nmsu.edu (GRUBB)  
Subject: Re: IDE vs SCSI  
Organization: New Mexico State University, Las Cruces, NM  
Lines: 44  
Distribution: world  
NNTP-Posting-Host: dante.nmsu.edu

DXB132@psuvm.psu.edu writes:

>In article <1qlbrlINN7rk@dns1.NMSU.Edu>, bgrubb@dante.nmsu.edu (GRUBB) says:  
>>In PC Magazine April 27, 1993:29 "Although SCSI is twice as fasst as ESDI,  
>>20% faster than IDE, and support up to 7 devices its acceptance ...has  
>>long been stalled by incompatability problems and installation headaches."

>I love it when magazine writers make stupid statements like that re:  
>performance. Where do they get those numbers? I'll list the actual  
>performance ranges, which should convince anyone that such a  
>statement is absurd:

>SCSI-I ranges from 0-5MB/s.

>SCSI-II ranges from 0-40MB/s.

>IDE ranges from 0-8.3MB/s.

>ESDI is always 1.25MB/s (although there are some non-standard versions)

ALL this shows is that YOU don't know much about SCSI.

SCSI-1 {with a SCSI-1 controller chip} range is indeed 0-5MB/s  
and that is ALL you have right about SCSI  
SCSI-1 {With a SCSI-2 controller chip}: 4-6MB/s with 10MB/s burst {8-bit}  
Note the INCREASE in SPEED, the Mac Quadra uses this version of SCSI-1  
so it DOES exist. Some PC use this set up too.  
SCSI-2 {8-bit/SCSI-1 mode}: 4-6MB/s with 10MB/s burst  
SCSI-2 {16-bit/wide or fast mode}: 8-12MB/s with 20MB/s burst  
SCSI-2 {32-bit/wide AND fast}: 15-20MB/s with 40MB/s burst

By your OWN data the "Although SCSI is twice as fast as ESDI" is correct  
With a SCSI-2 controller chip SCSI-1 can reach 10MB/s which is indeed  
"20% faster than IDE" {120% of 8.3 is 9.96}. ALL these SCSI facts have been  
posted to this newsgroup in my Mac & IBM info sheet {available by FTP on  
sumex-aim.stanford.edu (36.44.0.6) in the info-mac/report as  
mac-ibm-compare[version #].txt (It should be 173 but 161 may still be there)}

Part of this problem is both Mac and IBM PC are inconsiant about what SCSI  
is which. Though it is WELL documented that the Quadra has a SCSI-2 chip  
an Apple salesperson said "it uses a fast SCSI-1 chip" {Not at a 6MB/s,  
10MB/s burst it does not. SCSI-1 is 5MB/s maximum synchronous and Quadra  
uses ANsynchronous SCSI which is SLOWER} It seems that Mac and IBM see  
SCSI-1 interface and think 'SCSI-1' when it maybe a SCSI-1 interface driven  
in the machine by a SCSI-2 controller chip in 8-bit mode {Which is MUCH  
FASTER then true SCSI-1 can go}.

Don't slam an article because you don't understand what is going on.  
One reference for the Quadra's SCSI-2 controller chip is  
(Digital Review, Oct 21, 1991 v8 n33 p8(1)).

FILE NO:9

From: holmes7000@iscsvax.uni.edu  
Subject: WIn 3.0 ICON HELP PLEASE!  
Organization: University of Northern Iowa  
Lines: 10

I have win 3.0 and downloaded several icons and BMP's but I can't figure out  
how to change the "wallpaper" or use the icons. Any help would be appreciated.

Thanx,

-Brando

PS Please E-mail me

FILE NO:10

From: kerr@ux1.cso.uiuc.edu (Stan Kerr)  
Subject: Re: Sigma Designs Double up??  
Article-I.D.: ux1.C52u8x.B62  
Organization: University of Illinois at Urbana  
Lines: 29

jap10@po.CWRU.Edu (Joseph A. Pellettiere) writes:

> I am looking for any information about the Sigma Designs  
> double up board. All I can figure out is that it is a  
> hardware compression board that works with AutoDoubler, but  
> I am not sure about this. Also how much would one cost?

I've had the board for over a year, and it does work with Diskdoubler,  
but not with Autodoubler, due to a licensing problem with Stac Technologies,  
the owners of the board's compression technology. (I'm writing this  
from memory; I've lost the reference. Please correct me if I'm wrong.)

Using the board, I've had problems with file icons being lost, but it's  
hard to say whether it's the board's fault or something else; however,  
if I decompress the troubled file and recompress it without the board,  
the icon usually reappears. Because of the above mentioned licensing  
problem, the freeware expansion utility DD Expand will not decompress  
a board-compressed file unless you have the board installed.

Since Stac has its own product now, it seems unlikely that the holes in Autodoubler/Diskdoubler related to the board will be fixed. Which is sad, and makes me very reluctant to buy Stac's product since they're being so stinky. (But hey, that's competition.)

--

Stan Kerr

Computing & Communications Services Office, U of Illinois/Urbana  
Phone: 217-333-5217 Email: stankerr@uiuc.edu

## Considering only four Categories

In order to get faster execution times for this first example we will work on a partial dataset with only 4 categories out of the 20 available in the dataset:

In [3]:

```
categories = ['alt.atheism', 'soc.religion.christian', 'comp.graphics', 'sci.med']
twenty_train = fetch_20newsgroups(subset='train', categories=categories, shuffle=True, random_state=42)
print("\n Reduced Target Names:\n", twenty_train.target_names)
print("\n Reduced Target Length:\n", len(twenty_train.data))
print("\n First Document : ", twenty_train.data[0])
```

Reduced Target Names:  
['alt.atheism', 'comp.graphics', 'sci.med', 'soc.religion.christian']

Reduced Target Length:  
2257

First Document : From: sd345@city.ac.uk (Michael Collier)  
Subject: Converting images to HP LaserJet III?  
Nntp-Posting-Host: hampton  
Organization: The City University  
Lines: 14

Does anyone know of a good way (standard PC application/PD utility) to convert tif/img/tga files into LaserJet III format. We would also like to do the same, converting to HPGL (HP plotter) files.

Please email any response.

Is this the correct group?

Thanks in advance. Michael.

--

Michael Collier (Programmer)  
Email: M.P.Collier@uk.ac.city  
Tel: 071 477-8000 x3769  
Fax: 071 477-8565

The Computer Unit,  
The City University,  
London,  
EC1V 0HB.

## Extracting features from text files

**In order to perform machine learning on text documents, we first need to turn the text content into numerical feature vectors. The most intuitive way to do so is the bags of words representation:**

**Assign a fixed integer id to each word occurring in any document of the training set (for instance by building a dictionary from words to integer indices).**

**For each document  $i$ , count the number of occurrences of each word  $w$  and store it in  $X[i, j]$  as the value of feature  $j$  where  $j$  is the index of word  $w$  in the dictionary**

The bags of words representation implies that  $n_{\text{features}}$  is the number of distinct words in the corpus: this number is typically larger

than 100,000.

## Tokenizing text with scikit-learn:

Text preprocessing, tokenizing and filtering of stopwords are included in a high level component that is able to build a dictionary of features and transform documents to feature vectors:

## Word Occurrences

In [4]:

```
from sklearn.feature_extraction.text import CountVectorizer
count_vect = CountVectorizer()
X_train_counts = count_vect.fit_transform(twenty_train.data)
print("\n(Target Length , Distinct Words):", X_train_counts.shape)
print("\n Frequency of the word algorithm:", count_vect.vocabulary_.get('algorithm'))
```

(Target Length , Distinct Words): (2257, 35788)

Frequency of the word algorithm: 4690

## From occurrences to frequencies

**Term Frequencies :** Divide the number of occurrences of each word in a document by the total number of words in the document: these new features are called tf for Term Frequencies.

**Another refinement on top of tf is to downscale weights for words that occur in many documents in the corpus and are therefore less informative than those that occur only in a smaller portion of the corpus. This downscaling is called tf-idf for “Term Frequency times Inverse Document Frequency”.**

**Both tf and tf-idf can be computed as follows:**

In [5]:

```
from sklearn.feature_extraction.text import TfidfTransformer
tf_transformer = TfidfTransformer(use_idf=False).fit(X_train_counts)
X_train_tf = tf_transformer.transform(X_train_counts)
X_train_tf.shape
```

Out[5]:

(2257, 35788)

**In the above example-code, we firstly use the fit(..) method to fit our estimator to the data and secondly the transform(..) method to transform our count-matrix to a tf-idf representation.**

**These two steps can be combined to achieve the same end result faster by skipping redundant processing. This is done through using the fit\_transform(..) method as shown below, and as mentioned in the note in the previous section:**

In [6]:

```
tfidf_transformer = TfidfTransformer()
X_train_tfidf = tfidf_transformer.fit_transform(X_train_counts)
X_train_tfidf.shape
```

Out[6]:

(2257, 35788)



Now that we have our features, we can train a classifier to try to predict the category of a post. Let's start with a naive Bayes classifier, which provides a nice baseline for this task. scikit-learn includes several variants of this classifier; the one most suitable for word counts is the multinomial variant:

## Training a classifier

In [7]:

```
from sklearn.naive_bayes import MultinomialNB
clf = MultinomialNB().fit(X_train_tfidf, twenty_train.target)
```

## Predicting the Outcome

To try to predict the outcome on a new document we need to extract the features using almost the same feature extracting chain as before. The difference is that we call transform instead of fit\_transform on the transformers, since they have already been fit to the training set:

In [8]:

```
docs_new = ['God is love', 'OpenGL on the GPU is fast']
X_new_counts = count_vect.transform(docs_new)
X_new_tfidf = tfidf_transformer.transform(X_new_counts)

predicted = clf.predict(X_new_tfidf)

for doc, category in zip(docs_new, predicted):
    print('%r => %s' % (doc, twenty_train.target_names[category]))

'God is love' => soc.religion.christian
'OpenGL on the GPU is fast' => comp.graphics
```

## Building a pipeline

In order to make the vectorizer => transformer => classifier easier to work with, scikit-learn provides a Pipeline class that behaves like a compound classifier:

In [9]:

```
from sklearn.pipeline import Pipeline
text_clf = Pipeline([('vect', CountVectorizer()),
                     ('tfidf', TfidfTransformer()),
                     ('clf', MultinomialNB()),
                     ])
```

The names vect, tfidf and clf (classifier) are arbitrary. We shall see their use in the section on grid search, below. We can now train the model with a single command:

In [10]:

```
text_clf.fit(twenty_train.data, twenty_train.target)
```

Out[10]:

```
Pipeline(memory=None,
 steps=[('vect', CountVectorizer(analyzer='word', binary=False, decode_error='strict',
 dtype=<class 'numpy.int64'>, encoding='utf-8', input='content',
 lowercase=True, max_df=1.0, max_features=None, min_df=1,
 ngram_range=(1, 1), preprocessor=None, stop_words=None,
 strip...inear_tf=False, use_idf=True)), ('clf', MultinomialNB(alpha=1.0, class_prior=None,
 fit_prior=True))])
```

## Evaluation of the performance on the test set

In [11]:

```
#Evaluating the predictive accuracy of the model is equally easy:
import numpy as np
twenty_test = fetch_20newsgroups(subset='test',categories=categories, shuffle=True, random_state=42)
docs_test = twenty_test.data
predicted = text_clf.predict(docs_test)
np.mean(predicted == twenty_test.target)
```

Out[11]:

0.83488681757656458

scikit-learn further provides utilities for more detailed performance analysis of the results:

In [12]:

```
from sklearn import metrics
print(metrics.classification_report(twenty_test.target, predicted,
    target_names=twenty_test.target_names))
```

|                        | precision | recall | f1-score | support |
|------------------------|-----------|--------|----------|---------|
| alt.atheism            | 0.97      | 0.60   | 0.74     | 319     |
| comp.graphics          | 0.96      | 0.89   | 0.92     | 389     |
| sci.med                | 0.97      | 0.81   | 0.88     | 396     |
| soc.religion.christian | 0.65      | 0.99   | 0.78     | 398     |
| avg / total            | 0.88      | 0.83   | 0.84     | 1502    |

In [13]:

```
metrics.confusion_matrix(twenty_test.target, predicted)
```

Out[13]:

```
array([[192,  2,  6, 119],
       [ 2, 347,  4,  36],
       [ 2, 11, 322,  61],
       [ 2,  2,  1, 393]], dtype=int64)
```

**As expected the confusion matrix shows that posts from the newsgroups on atheism and christian are more often confused for one another than with computer graphics.**

In [ ]:

```
Reference : http://scikit-learn.org/stable/tutorial/text\_analytics/working\_with\_text\_data.html
```